

Model Development Phase

Date	11 July 2024
Team ID	SWTID1720013031
Project Title	Prediction and Analysis of Liver Patient Data Using Machine Learning
Maximum Marks	4 Marks

Initial Model Training Code, Model Validation and Evaluation Report

The initial model training code will be showcased in the future through a screenshot. The model validation and evaluation report will include classification reports, accuracy, and confusion matrices for multiple models, presented through respective screenshots.

Initial Model Training Code:

```
x = data.iloc[:, 0:-1]
y = data.iloc[:, -1]
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, random_state=0, test_size=0.3)
```

```
x_train.shape
```

```
(408, 10)
```

```
x_test.shape
```

```
(175, 10)
```

Logistic Regression

```
lr_param_grid = {  
    'penalty': ['l1', 'l2'],  
    'C': [0.001, 0.01, 0.1, 1.0, 10.0],  
    'solver': ['liblinear', 'saga']  
}
```

```
lr_s = LogisticRegression(max_iter=1000)
```

```
lr_grid_search = GridSearchCV(estimator=lr_s, param_grid=lr_param_grid, cv=5, scoring='accuracy', verbose=1, n_jobs=-1)
```

```
lr_grid_search.fit(x_train_final, y_train)
```

```
# Get best parameters
```

```
lr_best_params = lr_grid_search.best_params_
```

```
print("Best parameters for Logistic Regression:", lr_best_params)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

Best parameters for Logistic Regression: {'C': 1.0, 'penalty': 'l1', 'solver': 'liblinear'}

```
lr = LogisticRegression(**lr_best_params, max_iter=1000)
```

```
lr.fit(x_train_final, y_train)
```

```
▼ LogisticRegression  
LogisticRegression(max_iter=1000, penalty='l1', solver='liblinear')
```

```
y_pred_lr = lr.predict(x_test_final)
```

```
lr_acc = accuracy_score(y_pred_lr, y_test)
```

```
lr_acc
```

```
0.6971428571428572
```

```
print(classification_report(y_test, y_pred_lr))
```

	precision	recall	f1-score	support
1	0.72	0.93	0.81	122
2	0.50	0.15	0.23	53
accuracy			0.70	175
macro avg	0.61	0.54	0.52	175
weighted avg	0.65	0.70	0.64	175

```
lr_cross = cross_val_score(lr, x_train_final, y_train, scoring='accuracy', cv = 6)
lr_cross.mean()
```

```
0.7352941176470589
```

```
lr_cm = confusion_matrix(y_pred_lr, y_test)
lr_cm
```



```
array([[114, 45],
       [ 8, 8]], dtype=int64)
```

Support Vector Classifier (SVC)

```
svm_param_dist = {
    'C': uniform(0.1, 10),
    'kernel': ['linear', 'rbf'],
    'gamma': ['scale', 'auto']
}
```

```
svm_s = SVC()
```

```
# Initialize GridSearchCV
svm_random_search = RandomizedSearchCV(estimator=svm_s, param_distributions=svm_param_dist, n_iter=10, cv=3, scoring='accuracy', verbose=2, n_jobs=2,
# Fit GridSearchCV
svm_random_search.fit(x_train_final, y_train)
```

GridSearchCV

```
arch = RandomizedSearchCV(estimator=svm_s, param_distributions=svm_param_dist, n_iter=10, cv=3, scoring='accuracy', verbose=2, n_jobs=2, random_state=42)
rchCV
arch.fit(x_train_final, y_train)
```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

• RandomizedSearchCV

• estimator: SVC

• SVC

```
svm_best_params = svm_random_search.best_params_
print("Best parameters for SVC:", svm_best_params)
```

Best parameters for SVC: {'C': 7.41993941811405, 'gamma': 'scale', 'kernel': 'linear'}

```
svm = SVC(**svm_best_params)
```

```
svm.fit(x_train_final, y_train)
```

• SVC

SVC(C=7.41993941811405, kernel='linear')

```
y_pred_svm = svm.predict(x_test_final)
```

```
svm_acc = accuracy_score(y_pred_svm, y_test)
svm_acc
```

0.6971428571428572

```
print(classification_report(y_test, y_pred_svm))
```

	precision	recall	f1-score	support
1	0.70	1.00	0.82	122
2	0.00	0.00	0.00	53
accuracy			0.70	175
macro avg	0.35	0.50	0.41	175
weighted avg	0.49	0.70	0.57	175

```
svm_cross = cross_val_score(svm, x_train_final, y_train, scoring='accuracy', cv = 6)
svm_cross.mean()
```

```
0.7181372549019608
```

```
svm_cm = confusion_matrix(y_pred_svm, y_test)
svm_cm
```

```
array([[122, 53],
       [ 0,  0]], dtype=int64)
```

Random Forest Classifier

```
rfc_param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}
```

```
rfc_s = RandomForestClassifier()
```

```
rfc_grid_search = GridSearchCV(estimator=rfc_s, param_grid=rfc_param_grid, cv=5, scoring='accuracy', verbose=1, n_jobs=-1)
rfc_grid_search.fit(x_train_final, y_train)
```

Fitting 5 folds for each of 108 candidates, totalling 540 fits

```
GridSearchCV
  estimator: RandomForestClassifier
    RandomForestClassifier
```

```
rfc_best_params = rfc_grid_search.best_params_
print("Best parameters for Random Forest Classifier:", rfc_best_params)
```

```
Best parameters for Random Forest Classifier: {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_split': 10, 'n_estimators': 100}
```

```
rfc = RandomForestClassifier(**rfc_best_params)
rfc.fit(x_train_final, y_train)
```

```
rfc = RandomForestClassifier(**rfc_best_params)
rfc.fit(x_train_final, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(max_depth=10, min_samples_leaf=2, min_samples_split=10)
```

```
ypred_rfc = rfc.predict(x_test_final)
```

```
rfc_acc = accuracy_score(ypred_rfc, y_test)
rfc_acc
```

```
0.6971428571428572
```

```
print(classification_report(y_test, ypred_rfc))
```

	precision	recall	f1-score	support
1	0.73	0.90	0.81	122
2	0.50	0.23	0.31	53
accuracy			0.70	175
macro avg	0.61	0.56	0.56	175
weighted avg	0.66	0.70	0.66	175

```
rfc_cross = cross_val_score(rfc, x_train_final, y_train, scoring='accuracy', cv = 6)
rfc_cross.mean()
```

```
0.7205882352941178
```

```
rfc_cm = confusion_matrix(ypred_rfc,y_test)
rfc_cm
```

```
array([[110, 41],
       [ 12, 12]], dtype=int64)
```

K Neighbors Classifier

```
knn_param_grid = {
    'n_neighbors': [3, 5, 7, 9],
    'weights': ['uniform', 'distance'],
    'metric': ['euclidean', 'manhattan']
}
```

```
knn_s = KNeighborsClassifier()
```

```
knn_grid_search = GridSearchCV(estimator=knn_s, param_grid=knn_param_grid, cv=5, scoring='accuracy', verbose=1, n_jobs=-1)
```

```
# Fit GridSearchCV
knn_grid_search.fit(x_train_final, y_train)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits

```
GridSearchCV
  estimator: KNeighborsClassifier
    KNeighborsClassifier
```

```
knn_best_params = knn_grid_search.best_params_
print("Best parameters for K Neighbors Classifier:", knn_best_params)
```

Best parameters for K Neighbors Classifier: {'metric': 'manhattan', 'n_neighbors': 3, 'weights': 'distance'}

```
knn = KNeighborsClassifier(**knn_best_params)
knn.fit(x_train_final, y_train)
```

```
KNeighborsClassifier
KNeighborsClassifier(metric='manhattan', n_neighbors=3, weights='distance')
```

```
ypred_knn = knn.predict(x_test_final)
```

```
knn_acc = accuracy_score(ypred_knn, y_test)
knn_acc
```

0.6628571428571428

```
print(classification_report(y_test, ypred_knn))
```

```

              precision    recall  f1-score   support

     1         0.74        0.80        0.77        122
     2         0.43        0.36        0.39         53

 accuracy          0.66        0.66        0.66        175
 macro avg         0.59        0.58        0.58        175
 weighted avg      0.65        0.66        0.65        175

```

```
knn_cross = cross_val_score(knn, x_train_final, y_train, scoring='accuracy', cv = 6)
knn_cross.mean()
```

```
0.7181372549019609
```

```
knn_cm = confusion_matrix(ypred_knn,y_test)
knn_cm
```

```
array([[97, 34],
       [25, 19]], dtype=int64)
```

```

pickle.dump(svm, open('svm_liver_analysis.pkl', 'wb'))
pickle.dump(rfc, open('rfc_liver_analysis.pkl', 'wb'))
pickle.dump(knn, open('knn_liver_analysis.pkl', 'wb'))
pickle.dump(lr, open('lr_liver_analysis.pkl', 'wb'))

```

Model Validation and Evaluation Report:

Model	Classification Report	Accuracy	Confusion Matrix

Logistic Regression	<pre>print(classification_report(y_test, y_pred_lr))</pre> <table><thead><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr></thead><tbody><tr><td>1</td><td>0.72</td><td>0.93</td><td>0.81</td><td>122</td></tr><tr><td>2</td><td>0.50</td><td>0.15</td><td>0.23</td><td>53</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.70</td><td>175</td></tr><tr><td>macro avg</td><td>0.61</td><td>0.54</td><td>0.52</td><td>175</td></tr><tr><td>weighted avg</td><td>0.65</td><td>0.70</td><td>0.64</td><td>175</td></tr></tbody></table>		precision	recall	f1-score	support	1	0.72	0.93	0.81	122	2	0.50	0.15	0.23	53	accuracy			0.70	175	macro avg	0.61	0.54	0.52	175	weighted avg	0.65	0.70	0.64	175	69.71% <pre>lr_acc = accuracy_score(y_pred_lr, y_test)</pre> <pre>lr_acc</pre> <pre>0.6971428571428572</pre>	<pre>lr_cm = confusion_matrix(y_pred_lr, y_test)</pre> <pre>lr_cm</pre> <pre>array([[114, 45], [8, 8]], dtype=int64)</pre>
	precision	recall	f1-score	support																													
1	0.72	0.93	0.81	122																													
2	0.50	0.15	0.23	53																													
accuracy			0.70	175																													
macro avg	0.61	0.54	0.52	175																													
weighted avg	0.65	0.70	0.64	175																													
Support Vector Classifier (SVC)	<pre>print(classification_report(y_test, y_pred_svm))</pre> <table><thead><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr></thead><tbody><tr><td>1</td><td>0.70</td><td>1.00</td><td>0.82</td><td>122</td></tr><tr><td>2</td><td>0.00</td><td>0.00</td><td>0.00</td><td>53</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.70</td><td>175</td></tr><tr><td>macro avg</td><td>0.35</td><td>0.50</td><td>0.41</td><td>175</td></tr><tr><td>weighted avg</td><td>0.49</td><td>0.70</td><td>0.57</td><td>175</td></tr></tbody></table>		precision	recall	f1-score	support	1	0.70	1.00	0.82	122	2	0.00	0.00	0.00	53	accuracy			0.70	175	macro avg	0.35	0.50	0.41	175	weighted avg	0.49	0.70	0.57	175	69.71% <pre>svm_acc = accuracy_score(y_pred_svm, y_test)</pre> <pre>svm_acc</pre> <pre>0.6971428571428572</pre>	<pre>svm_cm = confusion_matrix(y_pred_svm, y_test)</pre> <pre>svm_cm</pre> <pre>array([[122, 53], [0, 0]], dtype=int64)</pre>
	precision	recall	f1-score	support																													
1	0.70	1.00	0.82	122																													
2	0.00	0.00	0.00	53																													
accuracy			0.70	175																													
macro avg	0.35	0.50	0.41	175																													
weighted avg	0.49	0.70	0.57	175																													
Random Forest Classifier	<pre>print(classification_report(y_test, ypred_rfc))</pre> <table><thead><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr></thead><tbody><tr><td>1</td><td>0.73</td><td>0.90</td><td>0.81</td><td>122</td></tr><tr><td>2</td><td>0.50</td><td>0.23</td><td>0.31</td><td>53</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.70</td><td>175</td></tr><tr><td>macro avg</td><td>0.61</td><td>0.56</td><td>0.56</td><td>175</td></tr><tr><td>weighted avg</td><td>0.66</td><td>0.70</td><td>0.66</td><td>175</td></tr></tbody></table>		precision	recall	f1-score	support	1	0.73	0.90	0.81	122	2	0.50	0.23	0.31	53	accuracy			0.70	175	macro avg	0.61	0.56	0.56	175	weighted avg	0.66	0.70	0.66	175	69.7% <pre>rfc_acc = accuracy_score(ypred_rfc, y_test)</pre> <pre>rfc_acc</pre> <pre>0.6971428571428572</pre>	<pre>rfc_cm = confusion_matrix(ypred_rfc, y_test)</pre> <pre>rfc_cm</pre> <pre>array([[110, 41], [12, 12]], dtype=int64)</pre>
	precision	recall	f1-score	support																													
1	0.73	0.90	0.81	122																													
2	0.50	0.23	0.31	53																													
accuracy			0.70	175																													
macro avg	0.61	0.56	0.56	175																													
weighted avg	0.66	0.70	0.66	175																													
K Neighbors Classifier	<pre>print(classification_report(y_test, ypred_knn))</pre> <table><thead><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr></thead><tbody><tr><td>1</td><td>0.74</td><td>0.80</td><td>0.77</td><td>122</td></tr><tr><td>2</td><td>0.43</td><td>0.36</td><td>0.39</td><td>53</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.66</td><td>175</td></tr><tr><td>macro avg</td><td>0.59</td><td>0.58</td><td>0.58</td><td>175</td></tr><tr><td>weighted avg</td><td>0.65</td><td>0.66</td><td>0.65</td><td>175</td></tr></tbody></table>		precision	recall	f1-score	support	1	0.74	0.80	0.77	122	2	0.43	0.36	0.39	53	accuracy			0.66	175	macro avg	0.59	0.58	0.58	175	weighted avg	0.65	0.66	0.65	175	66.2% <pre>knn_acc = accuracy_score(ypred_knn, y_test)</pre> <pre>knn_acc</pre> <pre>0.6628571428571428</pre>	<pre>knn_cm = confusion_matrix(ypred_knn, y_test)</pre> <pre>knn_cm</pre> <pre>array([[97, 34], [25, 19]], dtype=int64)</pre>
	precision	recall	f1-score	support																													
1	0.74	0.80	0.77	122																													
2	0.43	0.36	0.39	53																													
accuracy			0.66	175																													
macro avg	0.59	0.58	0.58	175																													
weighted avg	0.65	0.66	0.65	175																													