

Model Optimization and Tuning Phase

Date	11 July 2024
Team ID	SWTID1720013031
Project Title	Prediction and Analysis of Liver Patient Data Using Machine Learning
Maximum Marks	10 Marks

Model Optimization and Tuning Phase

The Model Optimization and Tuning Phase involves refining machine learning models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

Hyperparameter Tuning Documentation (6 Marks):

Model	Tuned Hyperparameters	Optimal Values
Logistic Regression	<p>Logistic Regression</p> <pre>lr_param_grid = { 'penalty': ['l1', 'l2'], 'C': [0.001, 0.01, 0.1, 1.0, 10.0], 'solver': ['liblinear', 'saga'] }</pre> <pre>lr_s = LogisticRegression(max_iter=1000)</pre>	<pre>lr_grid_search.fit(x_train_final, y_train) # Get best parameters lr_best_params = lr_grid_search.best_params_ print("Best parameters for Logistic Regression:", lr_best_params) Fitting 5 folds for each of 20 candidates, totalling 100 fits Best parameters for Logistic Regression: {'C': 1.0, 'penalty': 'l1', 'solver': 'liblinear'}</pre> <pre>lr_acc = accuracy_score(y_pred_lr, y_test) lr_acc</pre> <p>0.6971428571428572</p> <pre>LogisticRegression LogisticRegression(max_iter=1000, penalty='l1', solver='liblinear')</pre> <p>Best parameters for Logistic Regression: {'C': 1.0, 'penalty': 'l1', 'solver': 'liblinear'}</p>

<p>Support Vector Classifier (SVC)</p>	<h3>Support Vector Classifier (SVC)</h3> <pre> svm_param_dist = { 'C': uniform(0.1, 10), 'kernel': ['linear', 'rbf'], 'gamma': ['scale', 'auto'] } svm_s = SVC() </pre>	<pre> svm_best_params = svm_random_search.best_params_ print("Best parameters for SVC:", svm_best_params) Best parameters for SVC: {'C': 7.41993941811405, 'gamma': 'scale', 'kernel': 'linear'} svm_acc = accuracy_score(y_pred_svm, y_test) svm_acc 0.6971428571428572 </pre> <div>SVC</div> <p>SVC(C=7.41993941811405, kernel='linear')</p> <p>Best parameters for SVC: {'C': 7.41993941811405, 'gamma': 'scale', 'kernel': 'linear'}</p>
<p>Random Forest Classifier</p>	<h3>Random Forest Classifier</h3> <pre> rfc_param_grid = { 'n_estimators': [100, 200, 300], 'max_depth': [None, 10, 20, 30], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4] } rfc_s = RandomForestClassifier() </pre>	<pre> rfc_best_params = rfc_grid_search.best_params_ print("Best parameters for Random Forest Classifier:", rfc_best_params) Best parameters for Random Forest Classifier: {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_split': 10, 'n_estimators': 100} rfc_acc = accuracy_score(ypred_rfc, y_test) rfc_acc 0.6971428571428572 </pre> <div>RandomForestClassifier</div> <p>RandomForestClassifier(max_depth=10, min_samples_leaf=2, min_samples_split=10)</p> <p>Best parameters for Random Forest Classifier: {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_split': 10, 'n_estimators': 100}</p>

<p>K Neighbors Classifier</p>	<h2 style="text-align: center;">K Neighbors Classifier</h2> <pre>knn_param_grid = { 'n_neighbors': [3, 5, 7, 9], 'weights': ['uniform', 'distance'], 'metric': ['euclidean', 'manhattan'] }</pre> <pre>knn_s = KNeighborsClassifier()</pre>	<pre>knn_best_params = knn_grid_search.best_params_ print("Best parameters for K Neighbors Classifier:", knn_best_params) Best parameters for K Neighbors Classifier: {'metric': 'manhattan', 'n_neighbors': 3, 'weights': 'distance'}</pre> <pre>knn_acc = accuracy_score(ypred_knn, y_test) knn_acc</pre> <p>0.6628571428571428</p> <div> <p>KNeighborsClassifier</p> <p>KNeighborsClassifier(metric='manhattan', n_neighbors=3, weights='distance')</p> </div> <p>Best parameters for K Neighbors Classifier:</p> <pre>{'metric': 'manhattan', 'n_neighbors': 3, 'weights': 'distance'}</pre>
---------------------------------------	---	---

Performance Metrics Comparison Report (2 Marks):

Model	Baseline Metric	Optimized Metric																														
Logistic Regression	<pre>lr = LogisticRegression(max_iter=1000) lr.fit(x_train_final, y_train)</pre>	<h2>Logistic Regression</h2> <pre>lr_param_grid = { 'penalty': ['l1', 'l2'], 'C': [0.001, 0.01, 0.1, 1.0, 10.0], 'solver': ['liblinear', 'saga'] }</pre> <pre>lr_s = LogisticRegression(max_iter=1000)</pre>																														
	<div>LogisticRegression</div> <div>LogisticRegression(max_iter=1000)</div>																															
	<pre>y_pred_lr = lr.predict(x_test_final)</pre> <pre>lr_acc = accuracy_score(y_pred_lr, y_test) lr_acc</pre>																															
	0.6857142857142857																															
	<pre>print(classification_report(y_test, y_pred_lr))</pre> <table><thead><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr></thead><tbody><tr><td>1</td><td>0.72</td><td>0.90</td><td>0.80</td><td>122</td></tr><tr><td>2</td><td>0.45</td><td>0.19</td><td>0.27</td><td>53</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.69</td><td>175</td></tr><tr><td>macro avg</td><td>0.59</td><td>0.55</td><td>0.53</td><td>175</td></tr><tr><td>weighted avg</td><td>0.64</td><td>0.69</td><td>0.64</td><td>175</td></tr></tbody></table>		precision	recall	f1-score	support	1	0.72	0.90	0.80	122	2	0.45	0.19	0.27	53	accuracy			0.69	175	macro avg	0.59	0.55	0.53	175	weighted avg	0.64	0.69	0.64	175	<div>LogisticRegression</div> <div>LogisticRegression(max_iter=1000, penalty='l1', solver='liblinear')</div>
	precision	recall	f1-score	support																												
1	0.72	0.90	0.80	122																												
2	0.45	0.19	0.27	53																												
accuracy			0.69	175																												
macro avg	0.59	0.55	0.53	175																												
weighted avg	0.64	0.69	0.64	175																												

	<pre>lr_cm = confusion_matrix(y_pred_lr,y_test) lr_cm array([[110, 43], [12, 10]], dtype=int64)</pre>	<pre>lr_acc = accuracy_score(y_pred_lr, y_test) lr_acc 0.6971428571428572 print(classification_report(y_test, y_pred_lr))</pre> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>1</td><td>0.72</td><td>0.93</td><td>0.81</td><td>122</td></tr><tr><td>2</td><td>0.50</td><td>0.15</td><td>0.23</td><td>53</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.70</td><td>175</td></tr><tr><td>macro avg</td><td>0.61</td><td>0.54</td><td>0.52</td><td>175</td></tr><tr><td>weighted avg</td><td>0.65</td><td>0.70</td><td>0.64</td><td>175</td></tr></table> <pre>lr_cm = confusion_matrix(y_pred_lr,y_test) lr_cm array([[114, 45], [8, 8]], dtype=int64)</pre>		precision	recall	f1-score	support	1	0.72	0.93	0.81	122	2	0.50	0.15	0.23	53	accuracy			0.70	175	macro avg	0.61	0.54	0.52	175	weighted avg	0.65	0.70	0.64	175																														
	precision	recall	f1-score	support																																																										
1	0.72	0.93	0.81	122																																																										
2	0.50	0.15	0.23	53																																																										
accuracy			0.70	175																																																										
macro avg	0.61	0.54	0.52	175																																																										
weighted avg	0.65	0.70	0.64	175																																																										
Support Vector Classifier (SVC)	<pre>svm = SVC() svm.fit(x_train_final, y_train)</pre> <div><div>▼ SVC</div><div>SVC()</div></div> <pre>y_pred_svm = svm.predict(x_test_final) svm_acc = accuracy_score(y_pred_svm, y_test) svm_acc 0.6971428571428572 print(classification_report(y_test, y_pred_svm))</pre> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>1</td><td>0.70</td><td>1.00</td><td>0.82</td><td>122</td></tr><tr><td>2</td><td>0.00</td><td>0.00</td><td>0.00</td><td>53</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.70</td><td>175</td></tr><tr><td>macro avg</td><td>0.35</td><td>0.50</td><td>0.41</td><td>175</td></tr><tr><td>weighted avg</td><td>0.49</td><td>0.70</td><td>0.57</td><td>175</td></tr></table> <pre>svm_cm = confusion_matrix(y_pred_svm,y_test) svm_cm array([[122, 53], [0, 0]], dtype=int64)</pre>		precision	recall	f1-score	support	1	0.70	1.00	0.82	122	2	0.00	0.00	0.00	53	accuracy			0.70	175	macro avg	0.35	0.50	0.41	175	weighted avg	0.49	0.70	0.57	175	<h3>Support Vector Classifier (SVC)</h3> <pre>svm_param_dist = { 'C': uniform(0.1, 10), 'kernel': ['linear', 'rbf'], 'gamma': ['scale', 'auto'] }</pre> <pre>svm = SVC(**svm_best_params) svm.fit(x_train_final, y_train)</pre> <div><div>▼ SVC</div><div>SVC(C=7.41993941811405, kernel='linear')</div></div> <pre>y_pred_svm = svm.predict(x_test_final) svm_acc = accuracy_score(y_pred_svm, y_test) svm_acc 0.6971428571428572 print(classification_report(y_test, y_pred_svm))</pre> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>1</td><td>0.70</td><td>1.00</td><td>0.82</td><td>122</td></tr><tr><td>2</td><td>0.00</td><td>0.00</td><td>0.00</td><td>53</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.70</td><td>175</td></tr><tr><td>macro avg</td><td>0.35</td><td>0.50</td><td>0.41</td><td>175</td></tr><tr><td>weighted avg</td><td>0.49</td><td>0.70</td><td>0.57</td><td>175</td></tr></table> <pre>svm_cm = confusion_matrix(y_pred_svm,y_test) svm_cm array([[122, 53], [0, 0]], dtype=int64)</pre>		precision	recall	f1-score	support	1	0.70	1.00	0.82	122	2	0.00	0.00	0.00	53	accuracy			0.70	175	macro avg	0.35	0.50	0.41	175	weighted avg	0.49	0.70	0.57	175
	precision	recall	f1-score	support																																																										
1	0.70	1.00	0.82	122																																																										
2	0.00	0.00	0.00	53																																																										
accuracy			0.70	175																																																										
macro avg	0.35	0.50	0.41	175																																																										
weighted avg	0.49	0.70	0.57	175																																																										
	precision	recall	f1-score	support																																																										
1	0.70	1.00	0.82	122																																																										
2	0.00	0.00	0.00	53																																																										
accuracy			0.70	175																																																										
macro avg	0.35	0.50	0.41	175																																																										
weighted avg	0.49	0.70	0.57	175																																																										

Random Forest Classifier

```
rfc = RandomForestClassifier()
rfc.fit(x_train_final, y_train)
```

▼ RandomForestClassifier
RandomForestClassifier()

```
ypred_rfc = rfc.predict(x_test_final)
```

```
rfc_acc = accuracy_score(ypred_rfc, y_test)
rfc_acc
```

0.6857142857142857

```
print(classification_report(y_test, ypred_rfc))
```

	precision	recall	f1-score	support
1	0.74	0.85	0.79	122
2	0.47	0.30	0.37	53
accuracy			0.69	175
macro avg	0.60	0.58	0.58	175
weighted avg	0.66	0.69	0.66	175

```
rfc_cm = confusion_matrix(ypred_rfc, y_test)
rfc_cm
```

```
array([[104, 37],
       [ 18, 16]], dtype=int64)
```

Random Forest Classifier

```
rfc_param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}
```

```
rfc_s = RandomForestClassifier()
```

▼ RandomForestClassifier
RandomForestClassifier(max_depth=10, min_samples_leaf=2, min_samples_split=10)

```
ypred_rfc = rfc.predict(x_test_final)
```

```
rfc_acc = accuracy_score(ypred_rfc, y_test)
rfc_acc
```

0.6971428571428572

```
print(classification_report(y_test, ypred_rfc))
```

	precision	recall	f1-score	support
1	0.73	0.90	0.81	122
2	0.50	0.23	0.31	53
accuracy			0.70	175
macro avg	0.61	0.56	0.56	175
weighted avg	0.66	0.70	0.66	175

```
rfc_cm = confusion_matrix(ypred_rfc, y_test)
rfc_cm
```

```
array([[110, 41],
       [ 12, 12]], dtype=int64)
```

K Neighbors Classifier

```
knn = KNeighborsClassifier()
knn.fit(x_train_final, y_train)
```

▼ KNeighborsClassifier
KNeighborsClassifier()

```
ypred_knn = knn.predict(x_test_final)
```

K Neighbors Classifier

```
knn_param_grid = {
    'n_neighbors': [3, 5, 7, 9],
    'weights': ['uniform', 'distance'],
    'metric': ['euclidean', 'manhattan']
}
```

```
knn_s = KNeighborsClassifier()
```

▼ KNeighborsClassifier
KNeighborsClassifier(metric='manhattan', n_neighbors=3, weights='distance')

```
knn_acc = accuracy_score(ypred_knn, y_test)
knn_acc

0.6457142857142857

print(classification_report(y_test, ypred_knn))
```

	precision	recall	f1-score	support
1	0.72	0.80	0.76	122
2	0.38	0.28	0.33	53
accuracy			0.65	175
macro avg	0.55	0.54	0.54	175
weighted avg	0.62	0.65	0.63	175

```
knn_cm = confusion_matrix(ypred_knn,y_test)
knn_cm

array([[98, 38],
       [24, 15]], dtype=int64)
```

```
ypred_knn = knn.predict(x_test_final)

knn_acc = accuracy_score(ypred_knn, y_test)
knn_acc

0.6628571428571428

print(classification_report(y_test, ypred_knn))
```

	precision	recall	f1-score	support
1	0.74	0.80	0.77	122
2	0.43	0.36	0.39	53
accuracy			0.66	175
macro avg	0.59	0.58	0.58	175
weighted avg	0.65	0.66	0.65	175

```
knn_cm = confusion_matrix(ypred_knn,y_test)
knn_cm

array([[97, 34],
       [25, 19]], dtype=int64)
```

Final Model Selection Justification (2 Marks):

Final Model	Reasoning
Logistic Regression	Logistic Regression offers a balance of high interpretability, slightly better generalization performance, and computational efficiency, making it the best choice for this dataset given the similar accuracy and F1-scores among the models.